

Optimisation and Decision Models

Introduction



- 1 About this Course
- 2 What is an Optimisation Problem?
- 3 Optimisation Terminology
- 4 Optimisation is Hard (in General)

- 1 About this Course**
- 2 What is an Optimisation Problem?**
- 3 Optimisation Terminology**
- 4 Optimisation is Hard (in General)**

This course introduces the *theory and application of modern optimisation* from a *business analytics perspective*.

The course should enable you to:

- 1 formulate *decision problems* in management and analytics as *mathematical optimisation problems*
- 2 solve the resulting optimisation problems using state-of-the-art *optimisation software packages*
- 3 appreciate the *computational complexity* of various classes of optimisation problems
- 4 get a *strong foundation in optimisation* in order to tackle later MSc courses (e.g. Machine Learning)

This course introduces the *theory and application of modern optimisation* from a *business analytics perspective*.

In order to achieve those goals within 20 hours, we prefer:

- 1 *modelling techniques* over *optimisation algorithms*
- 2 *intuitive (often visual) explanations* over *lengthy proofs*
- 3 *transferrable concepts* over *isolated facts*

This course will be tough, but we are here to help!

PART ONE: Linear Optimisation

- formulating linear programs
- solving linear programs
- linear programming duality

1.5 hours tutorial

8 hours lecture

4 individual homeworks: count for 10% of the final grade

PART TWO: Discrete Optimisation

- formulating discrete programs
- logical constraints
- solving discrete programs

1.5 hours tutorial

8 hours lecture

4 quizzes: count for 10% of the final grade

PART THREE: Nonlinear Optimisation

- recognising convexity
- special classes of convex programs
- solving convex programs

1 hour tutorial

4 hours lecture

4 quizzes: count for 40% of the final grade

PART ONE: Linear Optimisation

- formulating linear programs
- solving linear programs
- linear programming duality

1.5 hours tutorial

8 hours lecture
 $\times 2h$

PART TWO: Discrete Optimisation

- formulating discrete programs
- logical constraints
- solving discrete programs

1.5 hours tutorial

8 hours lecture
 $\times 2h$

PART THREE: Nonlinear Optimisation

- recognising convexity
- special classes of convex programs
- solving convex programs

1 hour tutorial

4 hours lecture

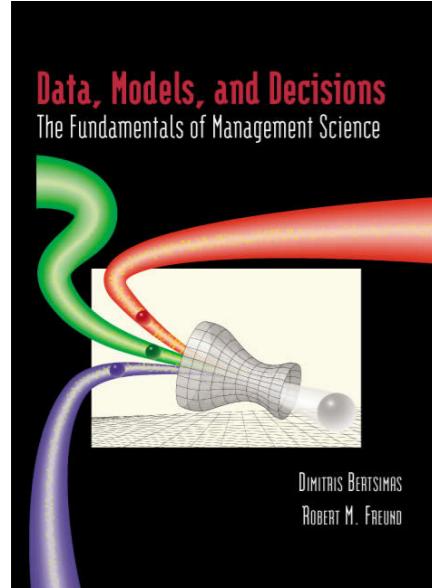
4 individual homeworks:

count for 10% of the final grade



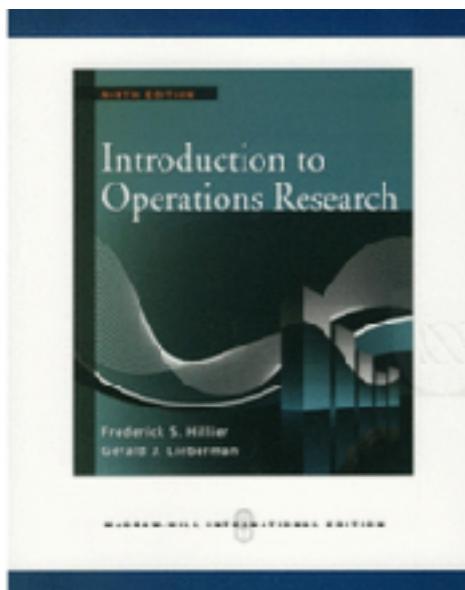
Timetabling Nightmare!

Mon 10	Tue 11	Wed 12	Thu 13	Fri 14	
		11:30-13:30 ACEX 554		9:30-11:30 LT 3	Lecture
				13:00-14:00 LT 2	Tutorial
Mon 17/24	Tue 18/25	Wed 19/26	Thu 20/27	Fri 21/28	Assignment
	16:00-17:00 ACEX 554	11:30-13:30 ACEX 554		11:30-12:30 LT 2	
				13:30-15:30 LT 2	
Mon 31	Tue 1	Wed 2	Thu 3	Fri 4	Tue 8
	16:00-17:00 ACEX 554	11:30-13:30 LT 3		11:30-12:30 LT 2	12:00-13:00 LT 1
				13:30-15:30 LT 2	



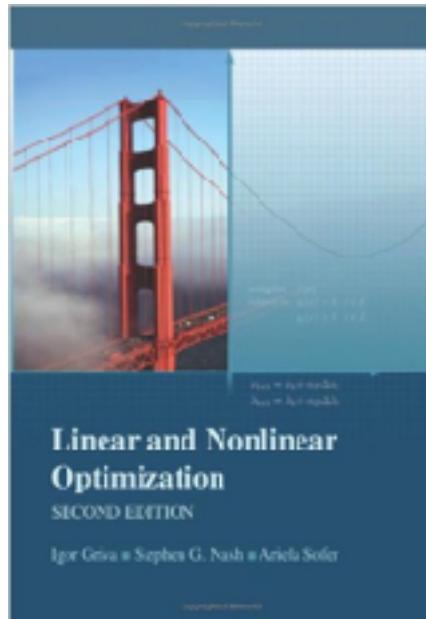
Bertsimas/Freund: *Data, Models, and Decisions*

- standard MBA textbook at MIT Sloan, HBS, Wharton, Columbia, Northwestern, ...
- good for basic concepts, exercises, MS Excel
- copies available at Central Library



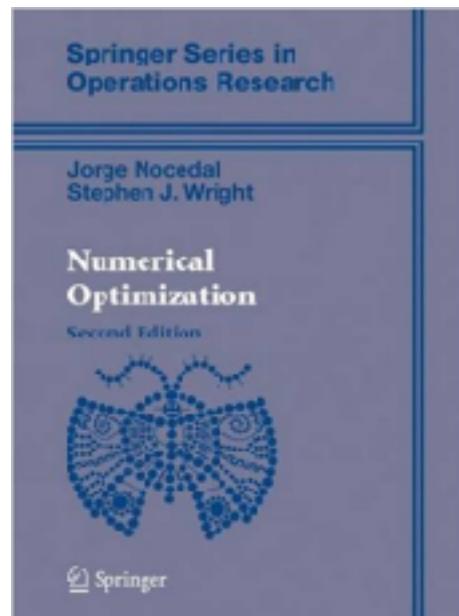
Hillier/Lieberman: *Introduction to OR*

- prize-winning Industrial Engineering classic (first version 1967), used throughout the world
- good coverage of duality + logical constraints
- copies available at Central Library



Griva/Nash/Sofe: *Linear and Nonlinear Opt.*

- good coverage of nonlinear optimisation
- goes beyond what is discussed in class
- 1 copy available at Central Library



Nocedal/Wright: *Numerical Optimization*

- classic textbook for nonlinear optimisation
- covers many details of real-life implementations
- goes beyond what is discussed in class
- copies available at Central Library

- 1 About this Course**
- 2 What is an Optimisation Problem?**
- 3 Optimisation Terminology**
- 4 Optimisation is Hard (in General)**

Leonhard Euler:^{*}

*“...nothing at all takes place
in the universe in which
some rule of maximum or
minimum does not appear...”*

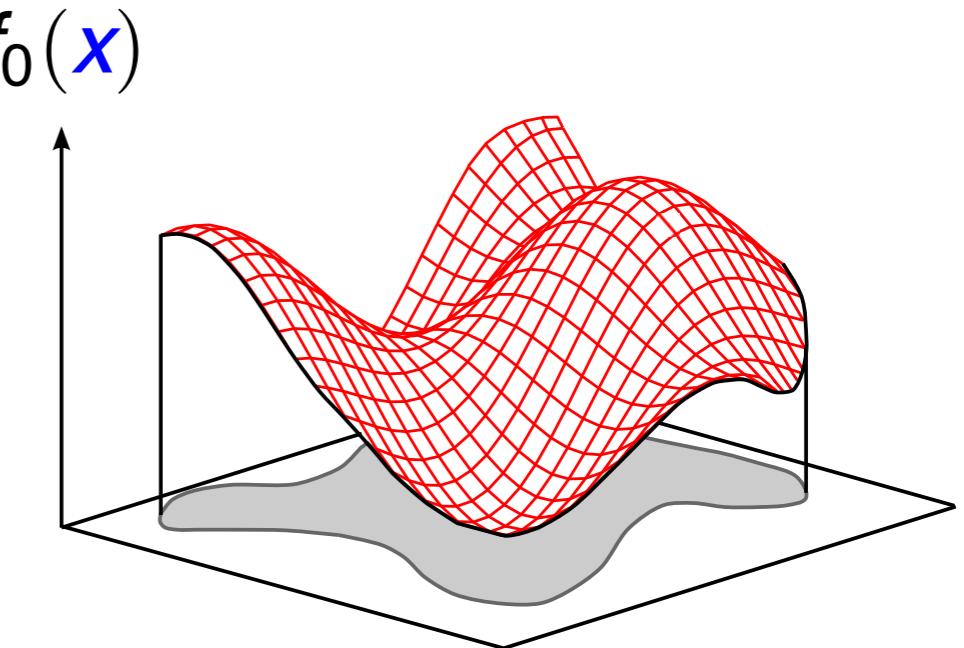


* Leonhard Euler. *Methodus inviendi lineas curvas maximi minimive proprietate gaudentes, sive, solutio problematis isoperimetrici latissimo sensu accepti*. Lausannæ; Genevæ: M.-M. Bousquet, 1744, p. 245.

The generic optimisation problem:

minimise $\underset{\mathbf{x}}{f_0(\mathbf{x})}$

subject to $\mathbf{x} \in \mathcal{X}$



where:

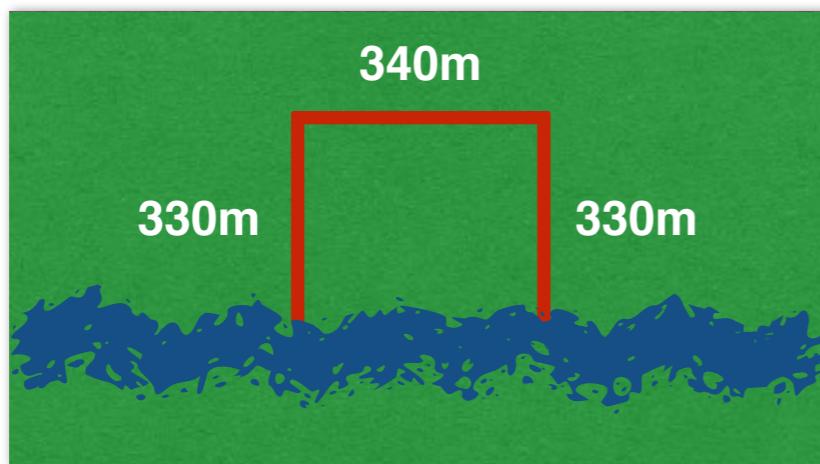
- $\mathbf{x} \in \mathbb{R}^n$ is the vector of *decision variables*
- $f_0 : \mathbb{R}^n \mapsto \mathbb{R}$ is the *objective function* (e.g. costs of \mathbf{x})
- $\mathcal{X} \subseteq \mathbb{R}^n$ is the *feasible region* (admissible values of \mathbf{x})

The feasible region is typically expressed through *constraints*:

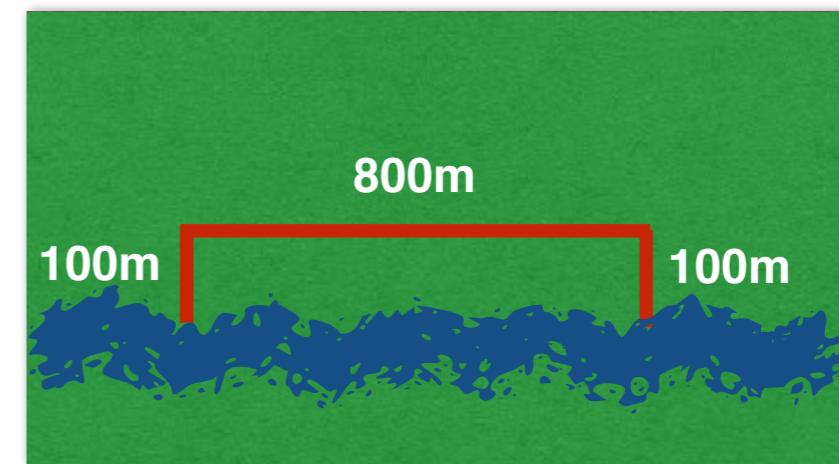
$$\begin{aligned} f_i(\mathbf{x}) &\leq 0 & \forall i = 1, \dots, m \\ h_i(\mathbf{x}) &= 0 & \forall i = 1, \dots, p \end{aligned}$$

where $f_i, h_i : \mathbb{R}^n \mapsto \mathbb{R}$.

A farmer has 1000m of fencing with which he wants to use to fence off a rectangular field that borders a straight river. Assuming that he does not need any fence along the river, what are the dimensions of the field that maximises the overall area?



area: $330\text{m} * 340\text{m} = 112,200\text{m}^2$



area: $100\text{m} * 800\text{m} = 80,000\text{m}^2$

maximise $w \cdot h$
subject to $w + 2h \leq 1000$
 $w, h \geq 0$

where:

- w is the field's width
- h is the field's height

* Adapted from https://cims.nyu.edu/~kiryl/Calculus/Section_4.5--Optimization%20Problems/Optimization_Problems.pdf.

Other Applications



Marriott International developed an optimisation-based real-time pricing strategy that improved hotel revenues by \$47m p.a.



Sloan Kettering Cancer Center improved prostate cancer therapeutics using optimisation, offering safer treatment and \$459m savings p.a.



TNT Express developed optimisation-based supply chain design and routing solutions that saved €207m over 4 years



dunnhumby uses optimisation to place, promote and price products for Tesco, P&G, Unilever, ...



Wien Energie uses optimisation to run a cascade of hydropower plants at Ottenstein Stausee (Austria)

- 1 About this Course
- 2 What is an Optimisation Problem?
- 3 **Optimisation Terminology**
- 4 Optimisation is Hard (in General)

$$\begin{array}{ll}\text{minimise}_{\mathbf{x}} & f_0(\mathbf{x}) \\ \text{subject to} & \mathbf{x} \in \mathcal{X}\end{array}$$

- \mathbf{x} is a *feasible solution* if $\mathbf{x} \in \mathcal{X}$, that is, if

$$\begin{array}{ll}f_i(\mathbf{x}) \leq 0 & \forall i = 1, \dots, m \\ h_i(\mathbf{x}) = 0 & \forall i = 1, \dots, p\end{array}$$

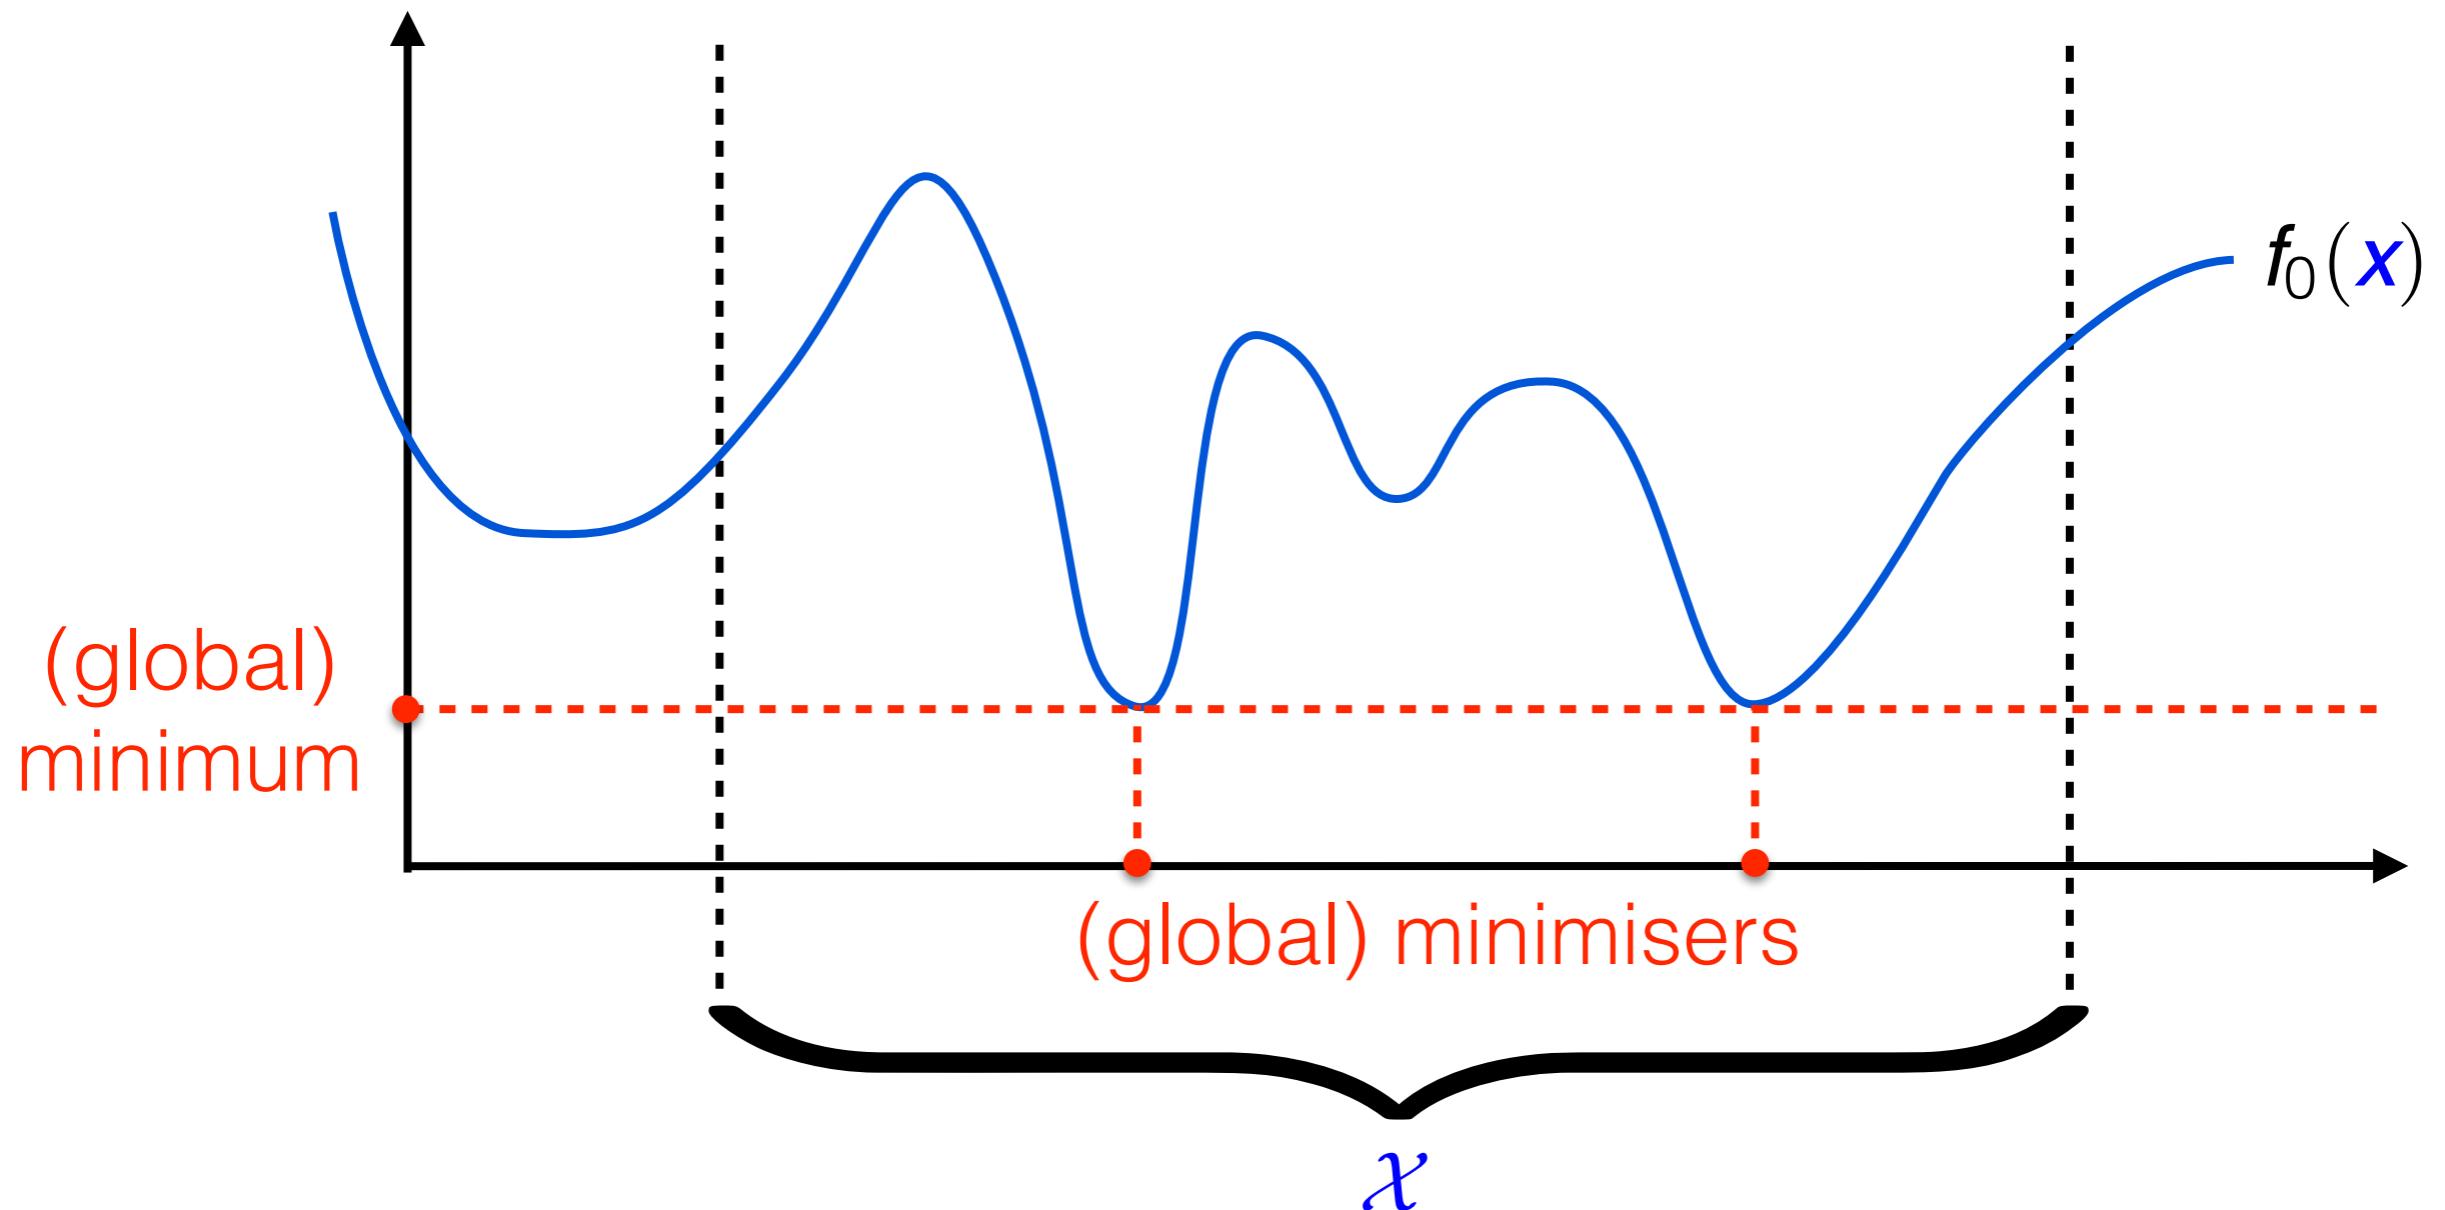
- \mathbf{x}^* is a *(global) minimiser* if \mathbf{x}^* is feasible and $f_0(\mathbf{x}^*) \leq f_0(\mathbf{x})$ for any feasible solution \mathbf{x} ; $f_0(\mathbf{x}^*)$ is the *(global) minimum*
- \mathbf{x}^* is a *local minimiser* if \mathbf{x}^* is feasible and $f_0(\mathbf{x}^*) \leq f_0(\mathbf{x})$ for any feasible solution \mathbf{x} “close to” \mathbf{x}^* ; $f_0(\mathbf{x}^*)$ is the *local minimum*. Formally put:

$\exists \delta > 0$ such that $f_0(\mathbf{x}^*) \leq f_0(\mathbf{x})$ for all $\mathbf{x} \in \mathcal{X}$ with $\|\mathbf{x} - \mathbf{x}^*\|_2 \leq \delta$

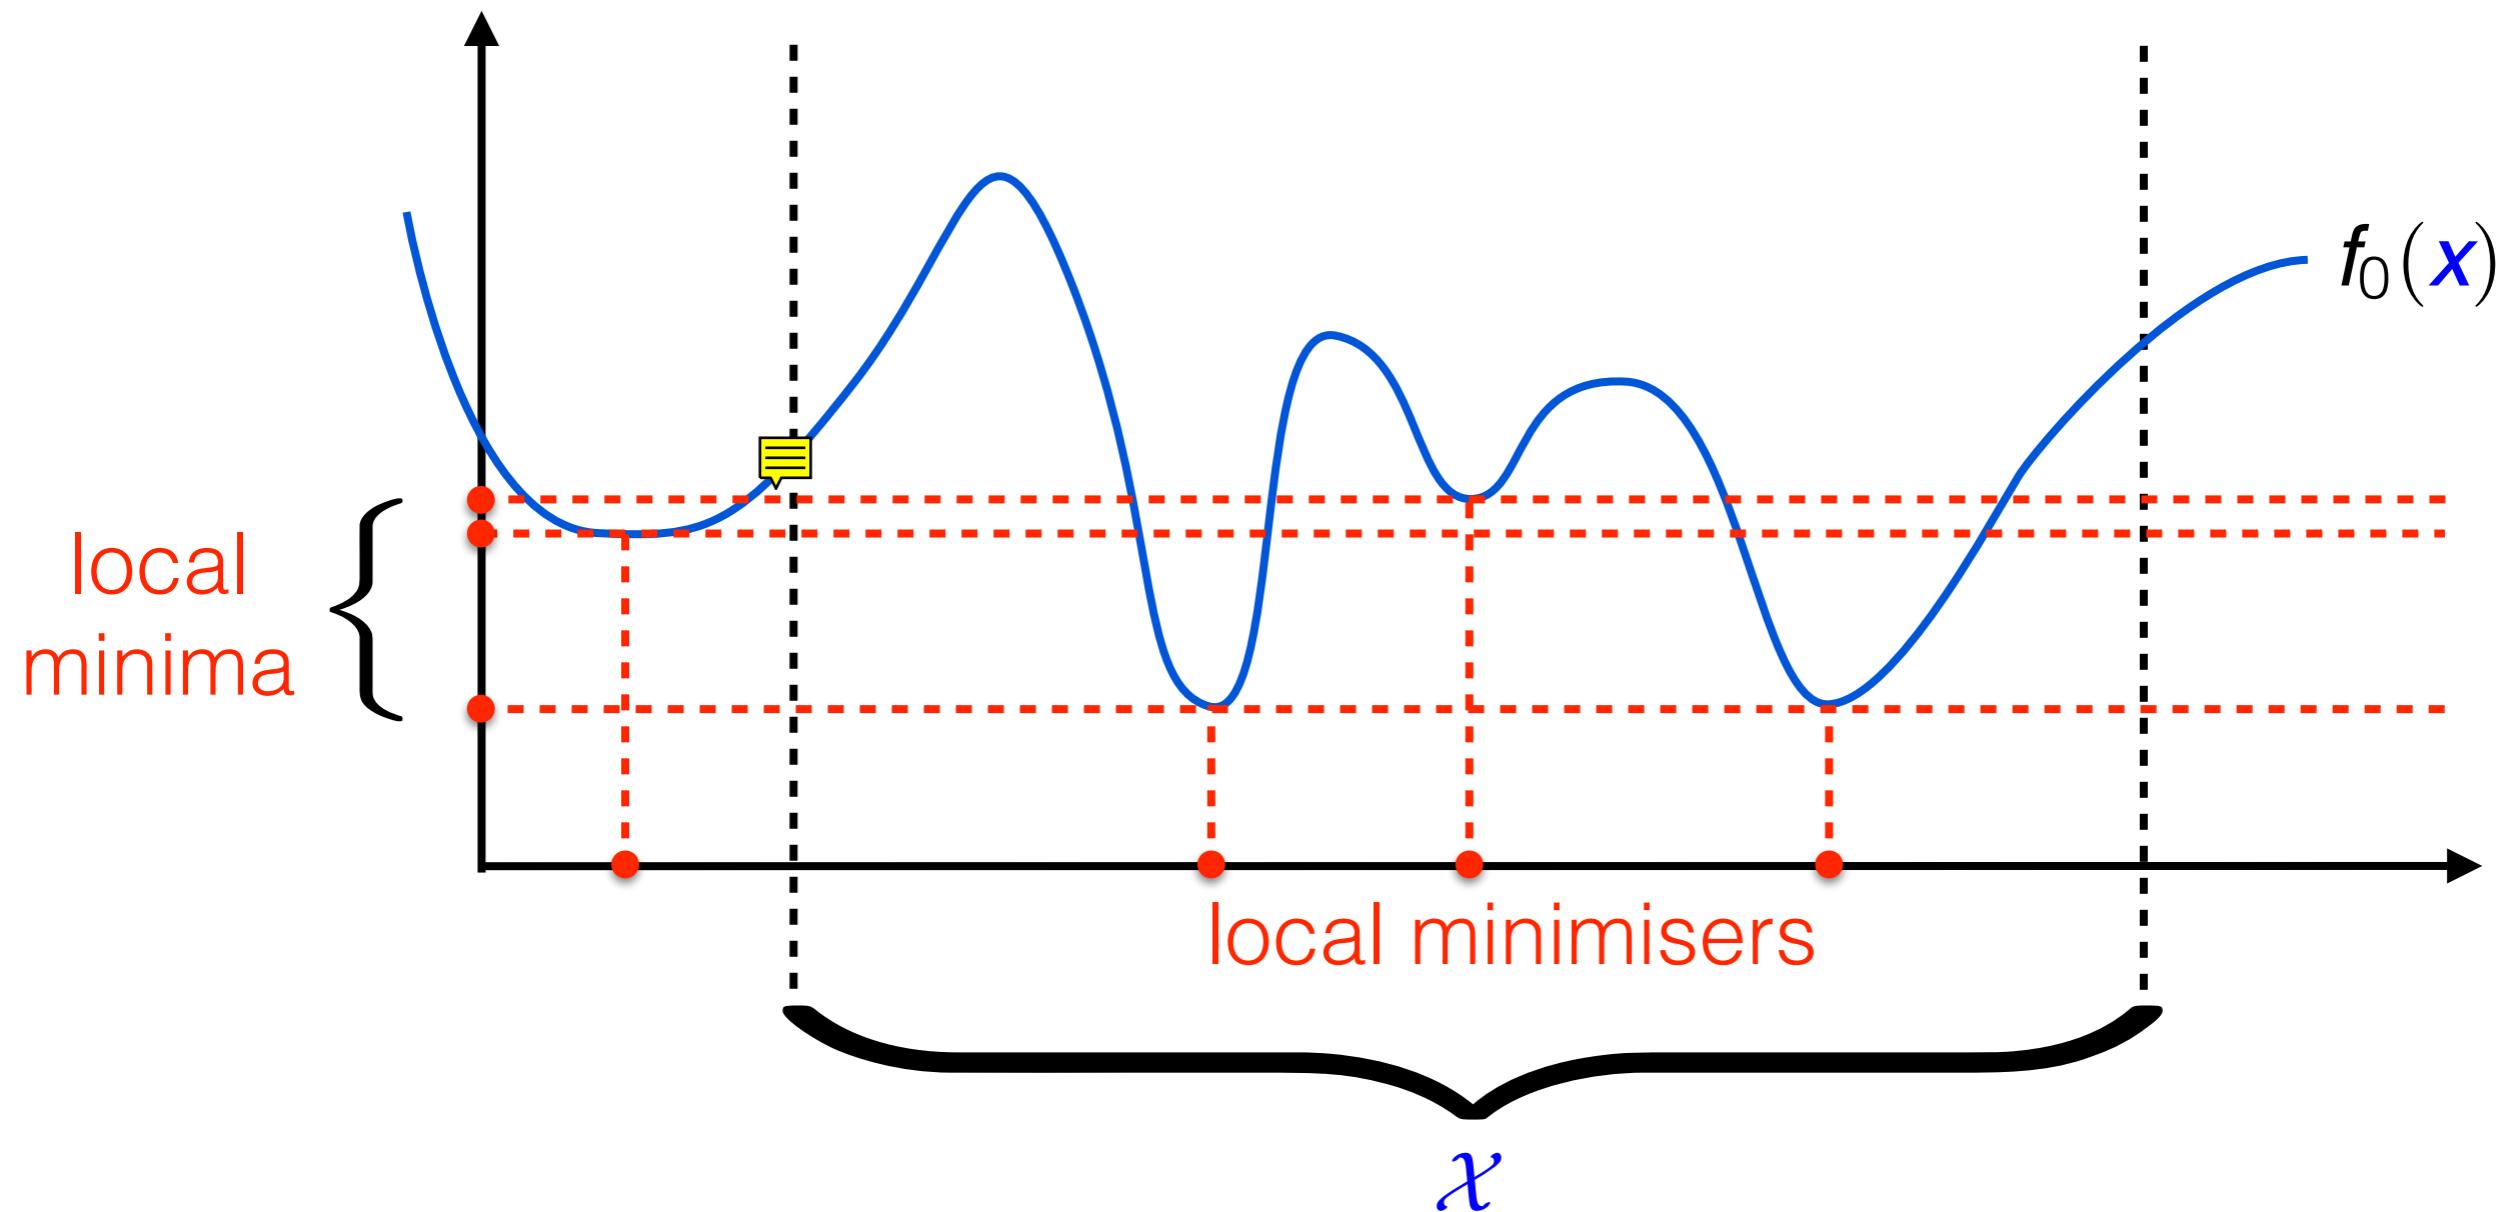
$$\begin{array}{ll}\text{minimise}_x & f_0(\mathbf{x}) \\ \text{subject to} & \mathbf{x} \in \mathcal{X}\end{array}$$

- for $\varepsilon > 0$, \mathbf{x}^* is called an *ε -(global) minimiser* if $\mathbf{x}^* \in \mathcal{X}$ and $f_0(\mathbf{x}^*) \leq f_0(\mathbf{x}) + \varepsilon$ for every feasible solution $\mathbf{x} \in \mathcal{X}$
- one can similarly define *ε -local minimisers*
- an optimisation problem is either: ☰
 - *infeasible* if the set \mathcal{X} is empty
 - *feasible and solvable* if it has at least one minimiser
 - *feasible but unbounded* if for every value θ there is a solution \mathbf{x} that satisfies $f(\mathbf{x}) \leq \theta$
 - *feasible and bounded but not solvable* if for every feas. sol. \mathbf{x} there is a feas. sol. \mathbf{x}' such that $f_0(\mathbf{x}') < f_0(\mathbf{x})$

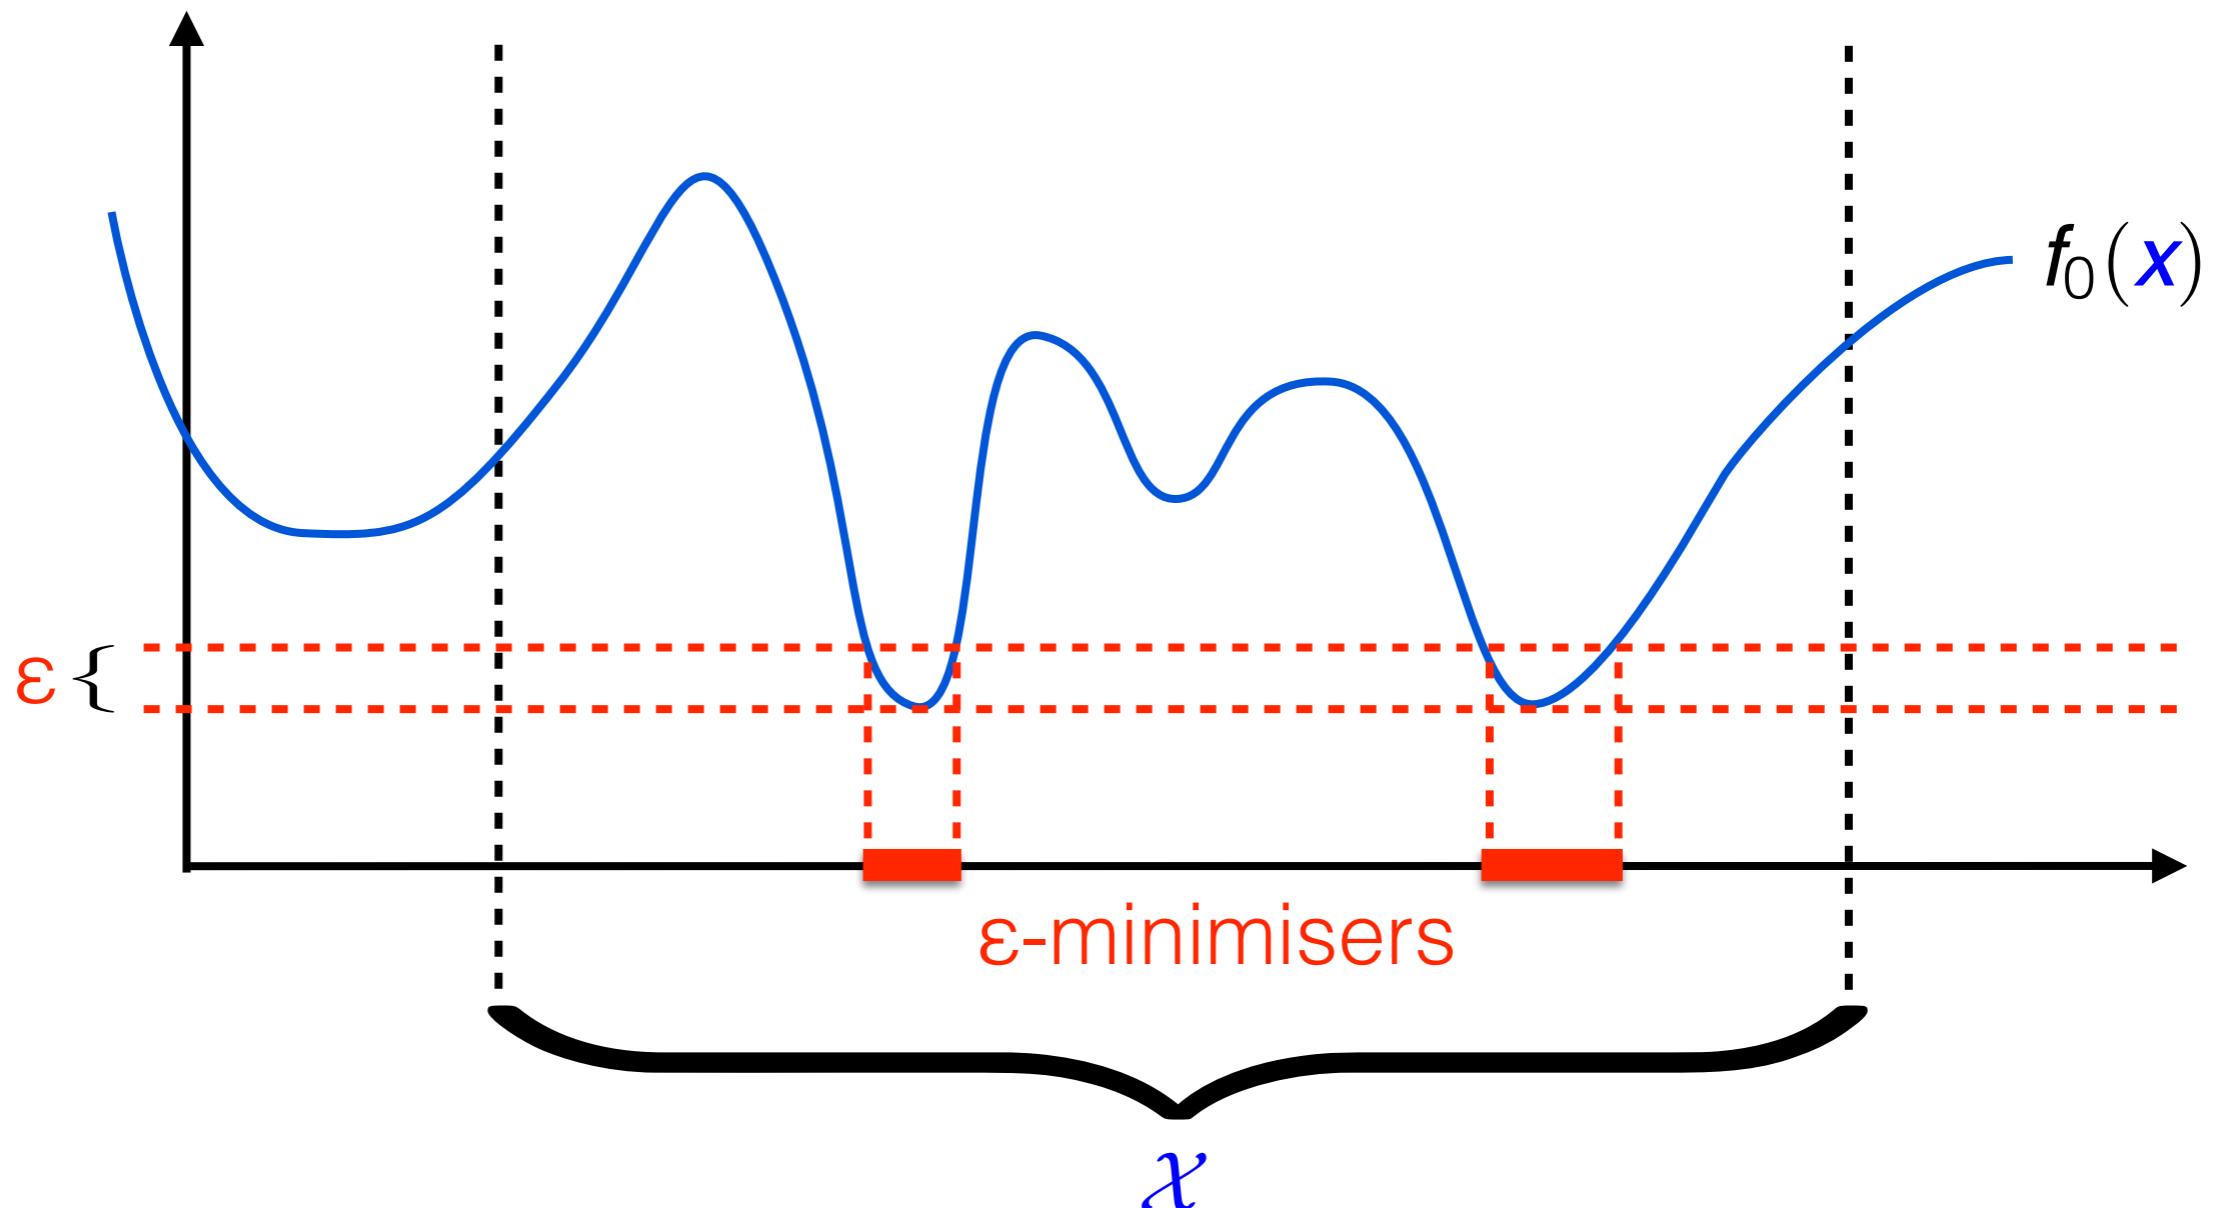
minimise $f_0(\mathbf{x})$
subject to $\mathbf{x} \in \mathcal{X}$



minimise $f_0(\mathbf{x})$
subject to $\mathbf{x} \in \mathcal{X}$



minimise $f_0(\mathbf{x})$
subject to $\mathbf{x} \in \mathcal{X}$



- 1 About this Course
- 2 What is an Optimisation Problem?
- 3 Optimisation Terminology
- 4 **Optimisation is Hard (in General)**

How long (in terms of # of operations) does the fastest algorithm take to solve a “simple” optimisation problem of the form

$$\begin{aligned} & \underset{\mathbf{x}}{\text{minimise}} && f_0(\mathbf{x}) \\ & \text{subject to} && \mathbf{x} \in [0, 1]^n \end{aligned}$$



where f_0 is Lipschitz-continuous with Lipschitz constant L ?¹

Answer:² To find an ε -minimiser for the problem, any algorithm requires at least $\left(\left\lfloor \frac{L}{2\varepsilon} \right\rfloor\right)^n$ evaluations of f_0 .

¹ A *Lipschitz-continuous function* f is a continuous function whose Lipschitz constant L describes how “rapidly” f changes ($L = 0$: f is flat; $L \rightarrow \infty$: f can change very rapidly).

² Proof: Y. Nesterov. *Introductory Lectures on Convex Optimization*. Kluwer, 2004. (Not required!)

How long (in terms of # of operations) does the fastest algorithm take to solve a “simple” optimisation problem of the form

$$\begin{aligned} & \underset{\mathbf{x}}{\text{minimise}} && f_0(\mathbf{x}) \\ & \text{subject to} && \mathbf{x} \in [0, 1]^n \end{aligned}$$

where f_0 is Lipschitz-continuous with Lipschitz constant L ?

Example: $L = 2$, $n = 10$, $\varepsilon = 0.01$ (tiny problem, low accuracy)

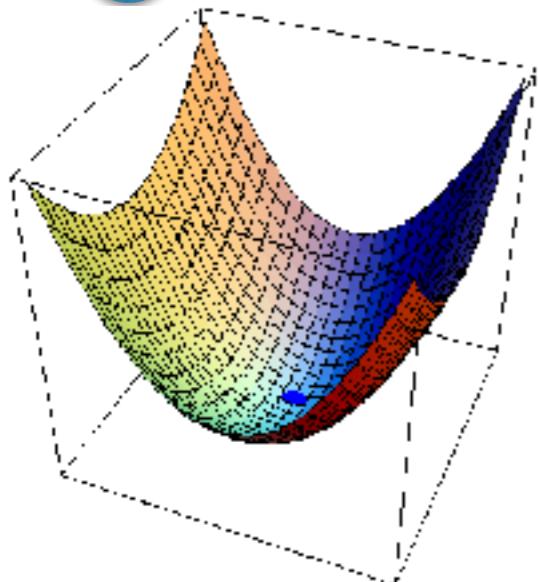
We need $\left(\left\lfloor \frac{2}{2 \cdot 0.01} \right\rfloor\right)^{10} = 10^{20}$ evaluations of f_0 .

If each evaluation requires 10 arithmetic operations and our computer can execute 10^6 arithmetic operations per second, then it takes **31,250,000 years!**

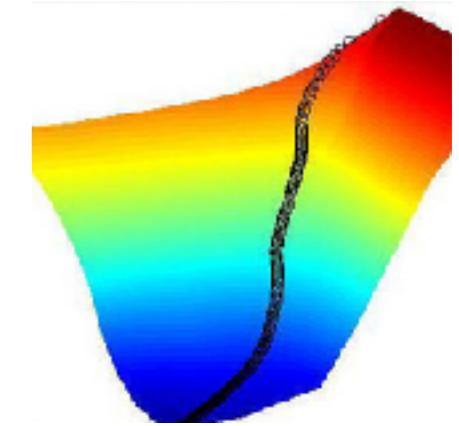


We can do much better, but for that we need:

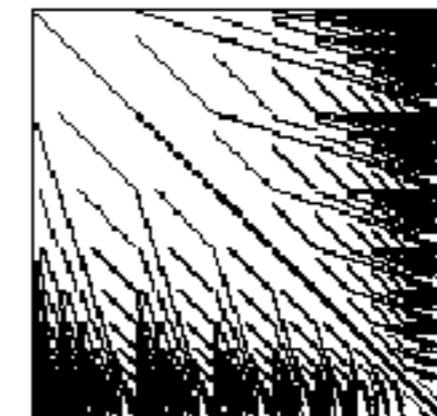
- 1 **Convexity:** implies that local minimum = global minimum



- 2 **Availability of derivatives:** helps us find a local minimum



- 3 **Sparsity:** each constraint involves only a few variables; makes underlying algebraic operations fast (e.g. matrix inverses)



*We can then reliably solve problems with **10,000,000 variables!***